

**POLITECNICO**  
**MILANO 1863**

# Presentation of OLA project

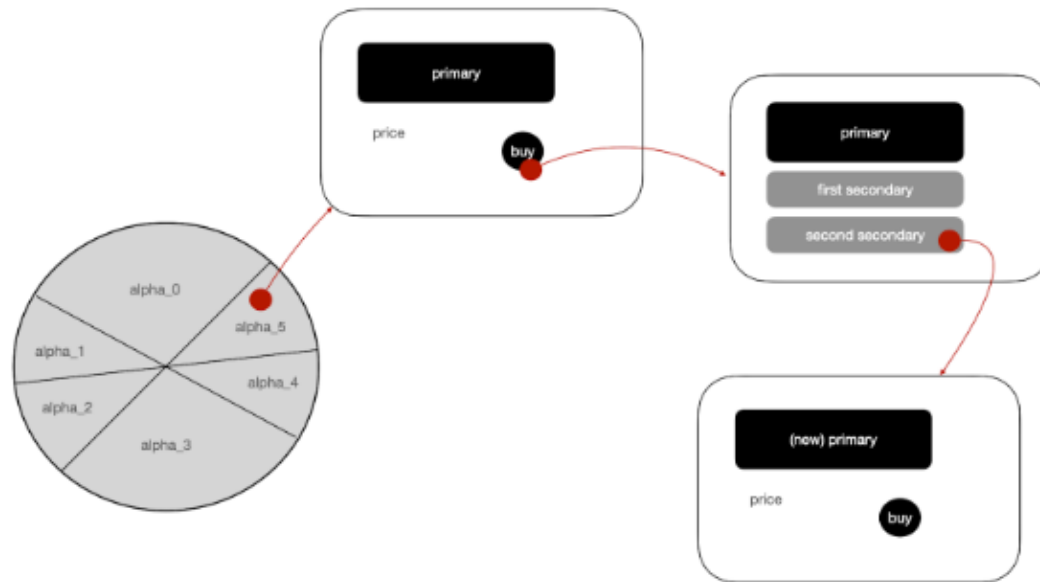
## **Pricing**

2021/2022

Alice Brugnoli, Marco Tonnarelli, Mattia Sabella, Andrea Isgrò, Giovanni Clini

# The e-commerce

Every day a random number of customers enters the e-commerce. Every single customer land on a webpage in which one of the 5 products is primary. If this product is added to the cart by the user, then two other products are recommended by the system in different slots. If the user clicks one of them, called secondary, then a new page is loaded in which the clicked one is displayed as primary together with its own price.




# STEP 1 - Environment

The actual environment is composed of the following classes:




There are also other two classes used in the scope of the environment:




# STEP 1 - Environment

## Environment and Contextual Environment



```
def round(self, configuration, prices, alpha_ratios, item_sold_mean):
    prices_configuration = np.zeros(self.n_prod)
    for product in range(self.n_prod):
        prices_configuration[product] = prices[product][configuration[product]]
    simulator = SocialInfluence(self.lambda_coeff, alpha_ratios, item_sold_mean,
        ↪ self.customer_class, prices_configuration, self.n_prod, configuration)
    simulator.simulation()
    return simulator.reward, simulator.bought, simulator.actual_users
```

## NonStationaryEnvironment



```
def round(self, configuration, prices, alpha_ratios, item_sold_mean, phase):
    prices_configuration = np.zeros(self.n_prod)
    for product in range(self.n_prod):
        prices_configuration[product] = prices[product][configuration[product]]
    simulator = SocialInfluence(self.lambda_coeff, alpha_ratios, item_sold_mean,
        ↪ self.customer_class, prices_configuration, self.n_prod, configuration)
    simulator.abrupt_simulation(phase)
    return simulator.reward, simulator.bought, simulator.actual_users
```

# STEP 1 - Customer Class

```
{
  "n_users": 100,
  "binary_features": [0, 0],
  "average_alphas": [0.02, 0.3, 0.01, 0.19, 0.3, 0.18],
  "graph_probabilities": [[0.0, 0.9, 0.01, 0.02, 0.07],
                          [0.25, 0.0, 0.25, 0.25, 0.25],
                          [0.5, 0.2, 0.0, 0.15, 0.15],
                          [0.8, 0.0, 0.1, 0.0, 0.1],
                          [0.1, 0.6, 0.1, 0.2, 0.0]],
  "conversion_rates": [[0.3, 0.4, 0.4, 0.3], [0.2, 0.3, 0.3, 0.3],
                       [0.4, 0.3, 0.2, 0.2], [0.2, 0.2, 0.5, 0.5],
                       [0.4, 0.2, 0.3, 0.4]],
  "reservation_prices": [[18, 15, 20, 30, 25]],
  "average_items_sold": [[3, 2, 5, 8], [2.1, 1.3, 0.6, 2.7],
                          [2, 3, 6, 8], [1, 7, 4, 4],
                          [0.7, 0.6, 1.4, 0.8]]
}
```



# STEP 1 – Social Influence

```
def simulation(self):
    # for each user, make a simulation
    for u in range(self.n_users):
        # assign initial product shown to the user,
        # given the dirichlet distribution
        initial_products = np.random.multinomial(1, self.dirichlet_probs)
        self.actual_users[np.where(initial_products == 1)[0]] += 1
        if initial_products[0] == 0:
            initial_products = initial_products[1:]
            # make a simulation
            self.graph_search(initial_products)
    # evaluate reward of the simulation
    self.evaluate_reward()

def graph_search(self, initial_active_nodes)

first_secondary_node = random.choices(range(len(prob_matrix[index, :])),
    ↪ prob_matrix[index, :], k=1)[0]
second_secondary_node = random.choices(range(len(prob_matrix[index, :])),
    ↪ prob_matrix[index, :], k=1)[0]
for i in range(5):
    if i == first_secondary_node:
        p_row[i] = p_row[first_secondary_node]
    elif i == second_secondary_node:
        p_row[i] = p_row[second_secondary_node] * self.lambda_coeff
    else:
        p_row[i] = 0.0
num_prod_clicked = np.random.randint(1, 3)
indx = random.choices(np.arange(0, 5), p_row, k=num_prod_clicked)
```

## STEP 2 – Optimization Problem

All the parameters are known and no online learner is required. The objective optimization problem is the maximization of the cumulative expected margin over all the products; in doing this, a greedy algorithm is used.

### CONFIGURATIONS:

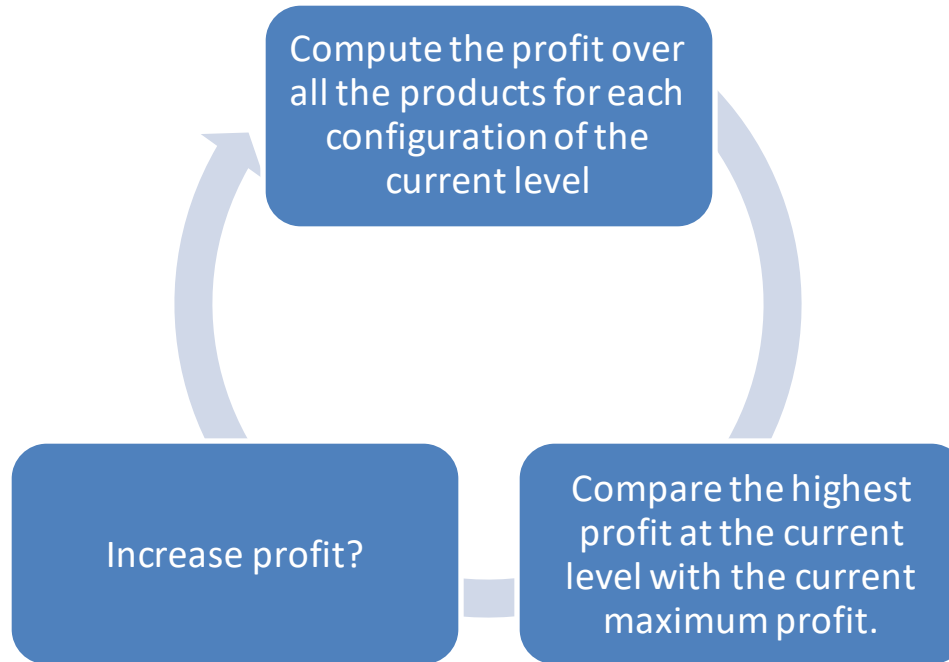
```
[array([16, 7, 5, 6, 4]), array([15, 14, 5, 6, 4]),  
array([15, 7, 10, 6, 4]), array([15, 7, 5, 12, 4]),  
array([15, 7, 5, 6, 8]), array([24, 14, 10, 12, 8]),  
array([16, 21, 10, 12, 8]), array([16, 14, 15, 12, 8]),  
array([16, 14, 10, 24, 8]), array([16, 14, 10, 12, 30]),  
array([32, 21, 15, 24, 30]), array([24, 28, 15, 24, 30]),  
array([24, 21, 40, 24, 30]), array([24, 21, 15, 26, 30]),  
array([24, 21, 15, 24, 50]), array([15, 7, 5, 6, 4])],
```

### INDEXES:

```
[array([1, 0, 0, 0, 0]), array([0, 1, 0, 0, 0]),  
array([0, 0, 1, 0, 0]), array([0, 0, 0, 1, 0]),  
array([0, 0, 0, 0, 1]), array([2, 1, 1, 1, 1]),  
array([1, 2, 1, 1, 1]), array([1, 1, 2, 1, 1]),  
array([1, 1, 1, 2, 1]), array([1, 1, 1, 1, 2]),  
array([3, 2, 2, 2, 2]), array([2, 3, 2, 2, 2]),  
array([2, 2, 3, 2, 2]), array([2, 2, 2, 3, 2]),  
array([2, 2, 2, 2, 3]), array([0, 0, 0, 0, 0])]
```

## STEP 2 – Optimization Algorithm

For every configuration level and for every user class:





# STEP 2 – Results

```
----- LEVEL 0 -----  
  
Customer class 0:  
4627.0  
Current marginal increase 92.59%  
The best configuration is number 3: [15 7 5 12 4]  
  
Customer class 1:  
1866.0  
Current marginal increase 158.35%  
The best configuration is number 3: [15 7 5 12 4]  
  
Customer class 2:  
789.0  
Current marginal increase 582.64%  
The best configuration is number 3: [15 7 5 12 4]
```

Figure 1: Algorithm Result (1)

```
----- LEVEL 1 -----  
  
Customer class 0:  
8911.0  
Current marginal increase 26.27%  
The best configuration is number 7: [16 14 15 12 8]  
  
Customer class 1:  
No better solution found  
Current marginal increase 0.00%  
The best configuration is number 3: [15 7 5 12 4]  
  
Customer class 2:  
No better solution found  
Current marginal increase 0.00%  
The best configuration is number 3: [15 7 5 12 4]
```

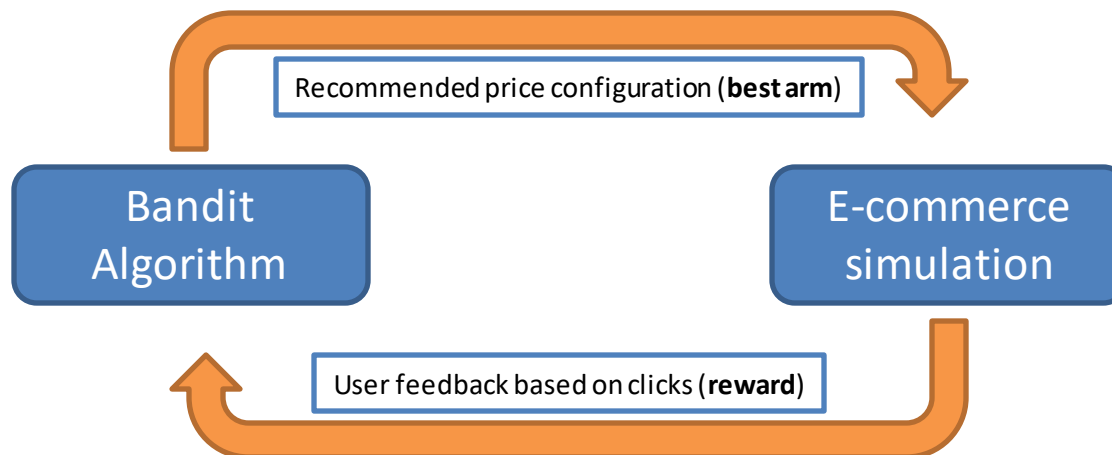
Figure 2: Algorithm Result (2)

```
----- LEVEL 2 -----  
  
Customer class 0:  
No better solution found  
Current marginal increase 0.00%  
The best configuration is number 7: [16 14 15 12 8]  
  
Customer class 1:  
No better solution found  
Current marginal increase 0.00%  
The best configuration is number 3: [15 7 5 12 4]  
  
Customer class 2:  
No better solution found  
Current marginal increase 0.00%  
The best configuration is number 3: [15 7 5 12 4]  
  
Process finished with exit code 0
```

Figure 3: Algorithm Result (3)

## STEP 3 - Optimzation with uncertain conversion rates

Step 3 provides the implementation of *TS* and *UCB* bandit algorithms in the scenario where *binary features are unknown*, and thus classes are to be considered aggregate, and the algorithms have to estimate the *conversion rates*.



The simulation of the e-commerce website has a duration of 200 days, and the number of experiments is 5.

In this step we also developed the clairvoyant estimator, which represents the best obtainable reward, and it is used to evaluate the regret.

## STEP 3 - Thompson Sampling and UCB

**Thompson Sampling:** it exploits prior knowledge to choose the best arm which maximizes the expected reward based on a randomly drawn belief.

Beta distribution is initialized uniformly to a matrix of ones. From this distribution, the conversion rates are evaluated for each configuration, and the one with the highest value is pulled. Then, simulation is run and the reward is collected. The update rule has been implemented as follows:

$$(\alpha_a, \beta_a) \leftarrow (\alpha_a, \beta_a) + (x_{a,t}, 1 - x_{a,t})$$

Python

```
self.beta_parameters[product, pulled_config[product], 0] =  
    ↪ self.beta_parameters[product, pulled_config[product], 0] +  
    ↪ bought[product]  
self.beta_parameters[product, pulled_config[product], 1] =  
    ↪ self.beta_parameters[product, pulled_config[product], 1] + total -  
    ↪ bought[product]
```

**UCB-1:** deterministic algorithm, where a confidence bound is associated to each arm. Confidence bound is updated according to the following formula:

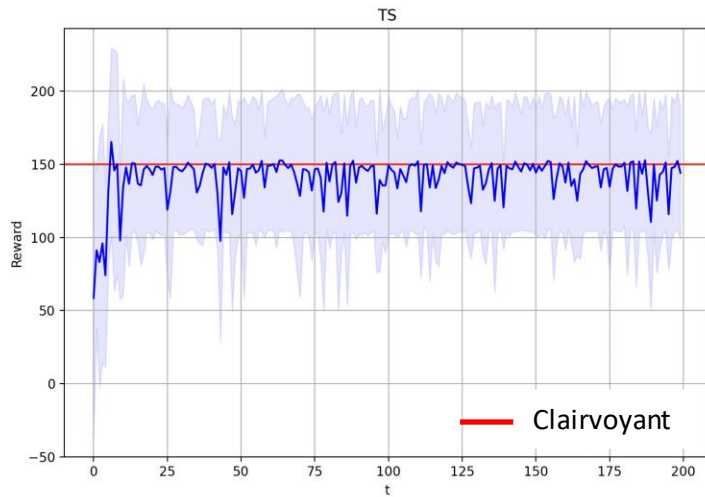
$$a_t \leftarrow \arg \max_{a \in A} \left\{ \bar{x}_a + \sqrt{\frac{2 \log(t)}{n_a(t-1)}} \right\}$$

Python


```
self.means[product, pulled_config[product]] = (self.means[product,  
    ↪ pulled_config[product]] * seen[product, pulled_config[product]] +  
    ↪ bought[product]) / tot_seen[product, pulled_config[product]]  
self.upper_bound[product, pulled_config[product]] = m.sqrt((2 *  
    ↪ m.log10(tot_samples[product])) / tot_seen[product,  
    ↪ pulled_config[product]])  
def pull_arm(self):  
    pulled_config_indexes = np.argmax(self.means + self.upper_bound, axis=1)  
    return pulled_config_indexes
```

Where the means are initialized to a matrix with all values equal to zero and the upper bound matrix to infinity.

## STEP 3 - Results: reward




TS reward



UCB reward

## STEP 3 - Results: regret



Regret comparison between UCB and TS algorithms

## STEP 4 - *Optimization with uncertain conversion rates, $\alpha$ ratios, and number of items sold per product*

This step reproduces the same conditions as the previous one. The only difference is that now also the alpha ratios and the number of products sold are unknown:

*$\alpha$  ratios:*

```
def estimate_alpha_ratios(old_starts, starts):  
    new_starts = old_starts + starts  
    return new_starts / np.sum(new_starts)
```

*Number (mean) of items sold per product:*

```
def estimate_items_for_each_product(mean, seen_since_day_before, unit_sold,  
    ↪ total_sold):  
    if mean * seen_since_day_before + unit_sold == 0:  
        return 0.  
    return (mean * seen_since_day_before + unit_sold) / total_sold
```



## STEP 4 - *Reward*




Figure 7: TS reward




Figure 8: UCB reward

## STEP 4 - *Cumulative regret*

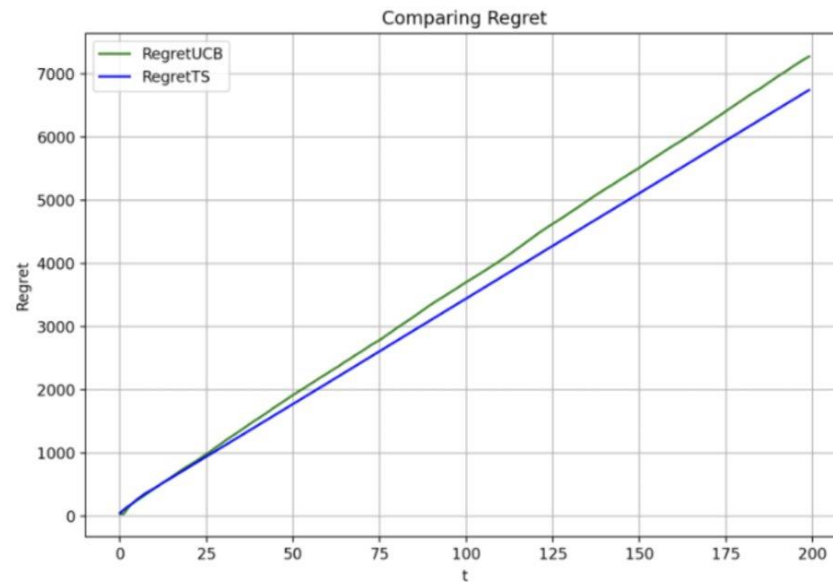


Figure 9: Cumulative regret

## STEP 5 – Graph Search

The function **Graph Search** is useful to simulate:

- How the products appear in the user page (e.g primary and secondary product shown).
- How the user interacts with the products.

The function is also relevant to collect data about:

- The number of clicks on the products.
- The number of units sold per product.
- The number of seen items per product.

## STEP 5 – Graph Search

At the beginning of the function we initialize the **Graph Probability Matrix**

```
def graph_search(self, initial_active_nodes):  
    # store the number of products, i.e., the nodes of the graph  
    prob_matrix = np.copy(self.graph_probs)  
    # initialize the history  
    history = []  
    # initialize the index of the activated node  
    index = 0  
    # initialize the history of the activated edges in the simulation  
    activated_edges_simulation = [False for _ in range(5)]
```

Then, we select the first and secondary product to be shown according to a **Probability Graph distribution** starting from the product selected by the user

```
while t < n_steps_max and len(order_of_parallel_product) != 0:  
    # index is the first product activated  
    index = np.where(order_of_parallel_product[0] == 1)[0][0]  
    # assign at random value given graph probabilities the first secondary node  
    first_secondary_node = random.choices(range(len(prob_matrix[index, :])), prob_matrix[index, :], k=1)[0]  
    # assign at random value given graph probabilities the second secondary node  
    second_secondary_node = random.choices(range(len(prob_matrix[index, :])), prob_matrix[index, :], k=1)[0]  
    # repeat assignment until the two products are different  
    while second_secondary_node == first_secondary_node:  
        second_secondary_node = random.choices(range(len(prob_matrix[index, :])), prob_matrix[index, :], k=1)[0]
```

## STEP 5 – Graph Search

Other essential step to decide the user choice on the next product clicked

```
for i in range(5):
    if i == first_secondary_node:
        p_row[i] = p_row[first_secondary_node]
    elif i == second_secondary_node:
        p_row[i] = p_row[second_secondary_node] * self.lambda_coeff
    else:
        p_row[i] = 0.0

# assign false to all the activated edges array to keep track of the one that will be selected (
# clicked secondary product)
activated_edges = [False for _ in range(5)]
# one or two secondary products are chosen num_prod_clicked = np.random.randint(1, 3) random choice
# of the index of the secondary product selected by the user (select one, but maybe could be two)
num_prod_clicked = np.random.randint(1, 3)
indx = random.choices(np.arange(0, 5), p_row, k=num_prod_clicked)
```

## STEP 5 – Graph Search

### Data Collection

```
# update the amount of unites of product purchased by the class of user
self.units_sold[i] += units_purchased
self.bought[i] += 1
self.customer_class.units_clicked_starting_from_a_primary[index][i] += 1
```

**Units\_clicked\_starting\_from\_a\_primary** is a matrix useful to keep track of the number of times a product is clicked starting from a given product (i.e defined by index).



## STEP 5 – Update Graph Probability




We keep the assumption of the previous steps, then we set the **Aggregate Graph Probability** to 1 for each cell of the matrix, except for the diagonal.

```
# NO AGGREGATE GRAPH BUT ONE GRAPH INITIALIZED TO 1
customer_class_aggregate.graph_probabilities = np.ones((5, 5))
np.fill_diagonal(customer_class_aggregate.graph_probabilities, 0)
```

After each simulation we update the values of the **Aggregate Graph Probability** for each product

```
for product in range(len(customer_class.units_clicked_starting_from_a_primary)):
    customer_class.graph_probabilities[product, :] = \
        (customer_class.graph_probabilities[product, :] * total_bought_day_before[product]
         + customer_class.units_clicked_starting_from_a_primary[product, :]) \
        / total_bought[product]
```

# STEP 5 – Results



# Step 6 - Non-stationary demand curve

Step 6 provides the implementation of two different *UCB-like algorithms*, which are able to tackle the same problem of step 3, but also with abrupt changes in demand curve.


In order to properly implement the *Change Detection algorithm* and the *Sliding Window* one, we developed a *NonStationaryEnvironment* module and the e-commerce simulation was modified to deal with changes. Also, the *clairvoyant* was modified such that it is able to evaluate the best reward for each phase.

The difference between the Change Detection algorithm and the Sliding Window one resides in the fact that the former detects the variation in the demand curve, while latter one has a fixed window size which flows along the timestamps of the experiment.

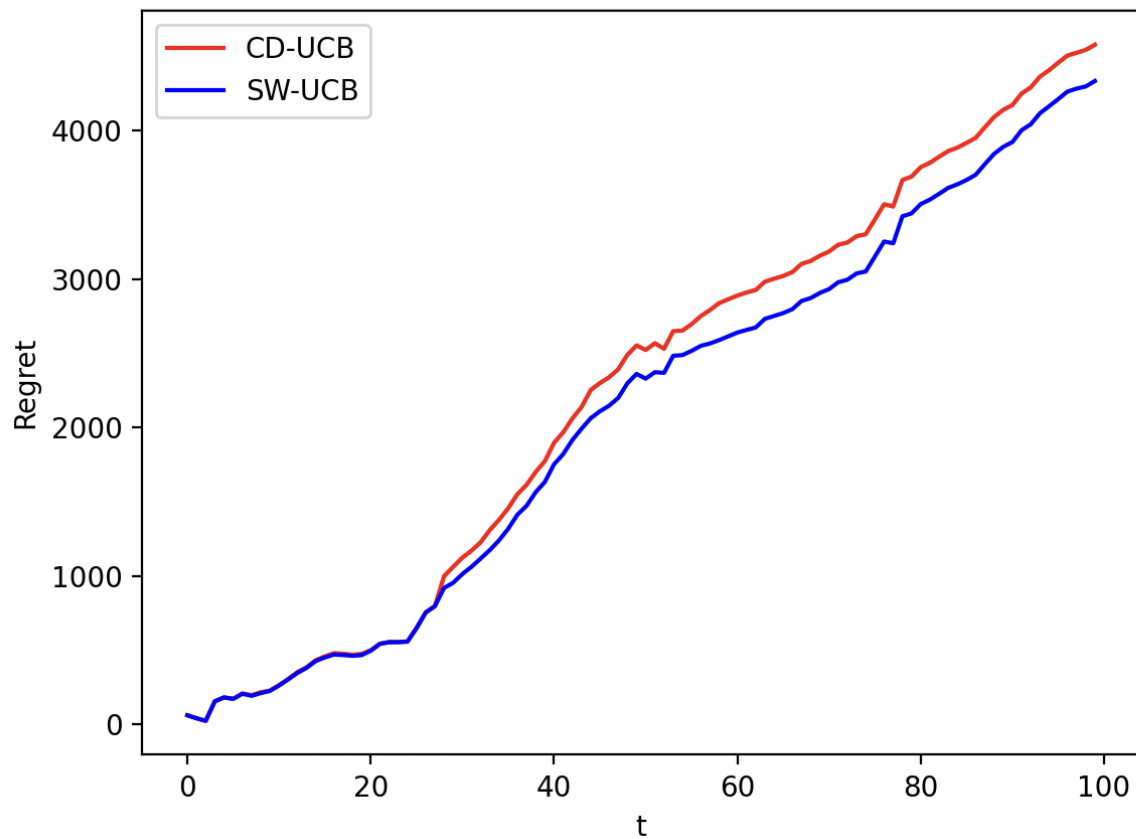
```
def detect_change(self, pulled_config):
    if self.t > 12:
        change_detected = False
        if not change_detected:
            for product in range(len(pulled_config)):
                last_mean =
                ↪ np.mean(self.conversion_rates[product][pulled_config[product]][:-self.size])
                new_mean =
                ↪ np.mean(self.conversion_rates[product][pulled_config[product]][-self.size:])
            if last_mean/new_mean >= self.delta or new_mean/last_mean >=
            ↪ self.delta:
                change_detected = True
                self.phase += 1
                self.means = np.zeros((self.n_products, self.n_arms))
                self.upper_bound = np.matrix(np.ones((self.n_products,
                ↪ self.n_arms)) * np.inf)
                self.phase_sizes.append(self.t)
            return True
```

Implementation of the change detection function

## Step 6 - Results: rewards




## Step 6 - Results: regret



## Step 7 - Context generation

The starting point of the following chapter is step number 4: conversion rates, alpha ratios and the number of items sold per product are unknown but this time the features can be observed by the e-commerce website. Every 14 days a Context Generation algorithm is executed to build a new features tree: we need to evaluate every possible partition of the space of the features, and for every one of these, we need to evaluate whether dis-aggregating such a subset is better than abstaining to do that. Practically, the features tree defines which context to use.





## Step 7 - Context generation

The algorithm is initialized with "no feature" classes which aggregate all the binary features. These classes - as the algorithm proceeds - are disaggregated and, after the two learners are trained on the initial partitions of data, a condition is evaluated and, if it is satisfied, the node is also disaggregated and the algorithm proceeds recursively. The splitting condition is as follows:

$$p_l \mu_l + p_r \mu_r \geq \mu$$

Python

```
def evaluate_split_condition(self, rewards0, rewards1, time):
    check = False
    l_lowerbound = self.lower_bound(rewards0, 5, 14)
    r_lowerbound = self.lower_bound(rewards1, 5, 14)
    p = self.assign_prob_context_occur(time)
    if p * (l_lowerbound + r_lowerbound) >= self.father_lower_bound:
        check = True
    return check, l_lowerbound, r_lowerbound
```

Where  $\mu$ ,  $\mu_l$  and  $\mu_r$  are the expected rewards,  $p_l$  and  $p_r$  are the probability weights and the lower bound  $\mu$  is:

$$\bar{x} - \sqrt{-\frac{\log(\delta)}{2|Z|}}$$

Python

```
def lower_bound(self, rewards, confidence, cardinality):
    return np.mean(rewards) - np.sqrt(-np.log10(confidence)/(cardinality * 2))
```

Where  $x$  is the expected reward computed by simulating the learner and  $|Z|$  is the size of the data.

## Step 7 - *Reward*




Figure 21: TS reward



Figure 22: UCB reward

## Step 7 – Cumulative Regret




Figure 23: Cumulative regret



**POLITECNICO**  
MILANO 1863